

# Differentiable Reproducibility Entropy via Kernel Density Estimation

## 1. The Problem

We have a system of cells that evolve into binary patterns (states of 0 or 1). We want to train a neural network to maximize the **reproducibility** of these patterns.

The standard metric is **Reproducibility Entropy**:

$$H = - \sum_{k \in \text{Unique Patterns}} p_k \log(p_k)$$

However, this metric is **non-differentiable** for two reasons:

1. **Thresholding:** The step function mapping  $s < 0.5 \rightarrow 0$  has zero gradient.
2. **Counting:** The `unique` function used to calculate probabilities ( $p_k$ ) is discrete and breaks the computation graph.

To fix this, we replace the discrete operations with differentiable approximations using **Soft Discretization** and **Kernel Density Estimation (KDE)**.

---

## 2. The Solution Pipeline

### Step A: Soft Discretization (The Sigmoid)

First, we smooth the "cliff" of the threshold function. Instead of a hard step, we use a sigmoid function with a temperature parameter  $T$ .

$$\tilde{s} = \sigma(T \cdot (s - 0.5)) = \frac{1}{1 + e^{-T(s-0.5)}}$$

- **Low  $T$ :** Smooth, linear-like transition.
- **High  $T$ :** Approaches a hard binary step function.

### Step B: Soft Matching (The Gaussian Kernel)

In the discrete world, two patterns  $x$  and  $y$  are either identical (distance 0) or not. In the continuous world, we measure similarity using a **Gaussian Kernel**.

We define the squared Euclidean distance between two patterns  $x_i$  and  $x_j$  as  $d_{ij}^2 = \|x_i - x_j\|^2$ .

The similarity (kernel value) is:

$$K(x_i, x_j) = \exp\left(-\frac{d_{ij}^2}{2\sigma^2}\right)$$

### Key Properties:

1. **If  $x_i \approx x_j$ :** Distance is  $\approx 0$ , so  $K \approx 1$ .
2. **If  $x_i \neq x_j$ :** Distance is large, so  $K \approx 0$ .
3. **No Normalization:** We deliberately do **not** include the standard normalization factor ( $\frac{1}{\sqrt{2\pi\sigma^2}}$ ). We want the "self-similarity" to be exactly 1.0, effectively counting "1 match".

## Step C: Soft Counting

To estimate how frequent a specific pattern  $x_i$  is within a batch of size  $N$ , we simply sum the similarities it has with every other pattern in the batch.

$$C_i = \sum_{j=1}^N K(x_i, x_j)$$

- **Interpretation:**  $C_i$  represents the "Soft Count" of pattern  $i$ . If there are 5 identical patterns,  $C_i \approx 5$ .

## Step D: Estimating Probabilities

We convert the soft count into an estimated probability  $\hat{p}_i$  by normalizing over the batch size  $N$ .

$$\hat{p}_i = \frac{C_i}{N}$$


---

## 3. The Entropy Calculation & "Double Counting"

We calculate the entropy using the **Resubstitution Estimator**, which averages the log-probability over the samples (rows) rather than summing over unique types.

$$H_{batch} = -\frac{1}{N} \sum_{i=1}^N \log(\hat{p}_i)$$

## Why does this work? (Addressing Double Counting)

It might seem incorrect to sum over  $N$  rows if many rows are identical. However, this summation naturally performs the weighting required by Shannon Entropy.

### The Proof:

Let there be a unique pattern type  $U$  that appears  $M$  times in the batch.

1. For any sample  $i$  that is an instance of  $U$ , the soft count is  $C_i \approx M$ .

2. The estimated probability is  $\hat{p}_i \approx \frac{M}{N}$ .
3. This specific value,  $\log(\frac{M}{N})$ , appears in our summation exactly  $M$  times (once for each row).

We can rewrite the sum by grouping the unique patterns:

$$\begin{aligned}
H_{batch} &= -\frac{1}{N} \sum_{i=1}^N \log(\hat{p}_i) \\
&= -\frac{1}{N} \left[ \sum_{U \in \text{Unique}} M_U \cdot \log \left( \frac{M_U}{N} \right) \right] \\
&= -\sum_{U \in \text{Unique}} \left( \frac{M_U}{N} \right) \cdot \log \left( \frac{M_U}{N} \right)
\end{aligned}$$

Since  $\frac{M_U}{N}$  is exactly the probability  $P(U)$ , we recover the standard Shannon Entropy formula:

$$H = -\sum_{U \in \text{Unique}} P(U) \log P(U)$$

Thus, iterating over the rows and "double counting" correctly applies the probability weight  $P(U)$  to the term.

---

## 4. JAX Implementation

Here is the implementation combining these concepts.

```

import jax
import jax.numpy as jnp

def differential_entropy_loss(final_states, temperature=10.0, sigma=0.5):
    """
    Calculates the differentiable entropy of the pattern distribution.

    Args:
        final_states: (Batch, Cells) array in range [0, 1]
        temperature: Steepness of the sigmoid (Soft Discretization)
        sigma: Bandwidth of the kernel (Soft Matching)

    Returns:
        Scalar entropy value.
        0.0      => Perfect reproducibility (All patterns identical)
        ln(Batch) => Zero reproducibility (All patterns unique)
    """
    batch_size = final_states.shape[0]

    # 1. Soft Discretization
    # Push values towards 0 or 1

```

```
soft_patterns = jax.nn.sigmoid(temperature * (final_states - 0.5))

# 2. Pairwise Distance Matrix
# Shape: (Batch, Batch, Cells)
diff = soft_patterns[:, None, :] - soft_patterns[None, :, :]

# Squared Euclidean Distance
# Shape: (Batch, Batch)
dists_sq = jnp.sum(diff**2, axis=-1)

# 3. Soft Counting (Unnormalized Gaussian Kernel)
#  $K(x, y) = \exp(-d^2 / 2\sigma^2)$ 
# We use log-space for stability:  $\log(K) = -d^2 / 2\sigma^2$ 
log_kernels = -dists_sq / (2 * sigma**2)

# 4. Log-Probability Estimation
#  $\log(p_i) = \log(\sum \exp(\log_kernels)) - \log(N)$ 
# This represents  $\log(\text{Count} / N)$ 
log_probs = jax.scipy.special.logsumexp(log_kernels, axis=1) -
jnp.log(batch_size)

# 5. Entropy Calculation (Resubstitution Estimator)
#  $H = -\text{Mean}(\log(p_i))$ 
entropy = -jnp.mean(log_probs)

# Optional: Normalize by system size if needed
# return entropy / final_states.shape[1]
return entropy
```